

Deep Learning Fundamentals

Perceptron, Neural Network, ANN & CNN — Complete Guide

Deep Learning Fundamentals

Perceptron, Neural Network, ANN & CNN — Complete Guide

1. PERCEPTRON

Simple Definition

A Perceptron is the most basic unit of a neural network. It takes some inputs, multiplies them by weights, adds a bias, and produces an output — like a tiny decision-maker that says "yes" or "no".

Technical Definition

A Perceptron is a linear binary classifier introduced by Frank Rosenblatt (1958). It computes a weighted sum of input features, adds a bias term, and passes the result through a step (activation) function to produce a binary output (0 or 1). It is the foundational building block of all modern neural networks.

Hinglish Explanation

Perceptron ek simple model hai jo multiple inputs leta hai ($x_1, x_2, x_3...$), unko weights se multiply karta hai, sabko add karta hai, aur ek bias add karke decide karta hai ki output 0 hoga ya 1. Isse hum bol sakte hain "single neuron ka mathematical model". Agar weighted sum ek threshold se zyada hai to output 1, warna 0.

Real-life Example


Suppose a bank wants to decide whether to approve a loan (1 = approve, 0 = reject) based on inputs like income, credit score, and existing debt. The perceptron assigns weights to each factor (credit score might matter most), sums them with a bias, and outputs a yes/no decision — similar to a simple rule-based approval system.

Why it is used

- It is the simplest building block to understand how neurons compute decisions.
- Used for linearly separable binary classification problems (e.g., AND/OR gate logic).
- Forms the conceptual base for understanding deeper networks (MLPs, ANN, CNN).
- Historically important — it sparked the entire field of neural networks.

Architecture Diagram

Perceptron Architecture



Step-by-step Working

- Take inputs $x_1, x_2, \dots, x_n$ (features of the data).
- Multiply each input by its corresponding weight: $w_1 \cdot x_1, w_2 \cdot x_2, \dots, w_n \cdot x_n$.
- Sum all weighted inputs and add a bias term b .
- Pass this sum (z) through an activation function (commonly a step function).
- If $z \geq \text{threshold}$, output = 1 (fire), else output = 0 (don't fire).
- During training, compare output with actual label, compute error, and update weights using the Perceptron Learning Rule.
- Repeat over many epochs until weights converge (for linearly separable data).

Mathematical Formula with Explanation

Weighted Sum (Net Input):

$$z = w_1 \cdot x_1 + w_2 \cdot x_2 + \dots + w_n \cdot x_n + b = (w \cdot x) + b$$

Activation (Step Function):

$$f(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases}$$

Weight Update Rule (Perceptron Learning Rule):

$$\begin{aligned} w_{\text{new}} &= w_{\text{old}} + \text{learning_rate} * (y_{\text{actual}} - y_{\text{predicted}}) * x \\ b_{\text{new}} &= b_{\text{old}} + \text{learning_rate} * (y_{\text{actual}} - y_{\text{predicted}}) \end{aligned}$$

Explanation: w is the weight vector, x is the input vector, b is bias, learning_rate controls step size of update, and $(y_{\text{actual}} - y_{\text{predicted}})$ is the error term. Weights move in the direction that reduces error.

Advantages & Disadvantages

Advantages:

- Simple to understand and implement.
- Fast computation, very few parameters.

- Guaranteed to converge if data is linearly separable (Perceptron Convergence Theorem).
- Forms the foundation for understanding complex networks.

Disadvantages:

- Cannot solve non-linearly separable problems (famous XOR problem).
- Only works for binary classification.
- Step function is not differentiable, so gradient-based optimization isn't directly possible.
- Limited representation power — only linear decision boundaries.

Interview Questions

- What is a perceptron and who invented it?
- Why can't a single perceptron solve the XOR problem?
- What is the difference between perceptron and sigmoid neuron?
- Explain the perceptron learning rule with formula.
- What is the Perceptron Convergence Theorem?
- Why is the step function not used in modern deep networks?
- How is a perceptron different from logistic regression?

Python Code

```
class Perceptron:
    def __init__(self, n_inputs, lr=0.1, epochs=50):
        self.weights = [0.0] * n_inputs
        self.bias = 0.0
        self.lr = lr
        self.epochs = epochs

    def activation(self, z):
        return 1 if z >= 0 else 0

    def predict(self, x):
        z = sum(w * xi for w, xi in zip(self.weights, x)) + self.bias
        return self.activation(z)

    def train(self, X, y):
        for _ in range(self.epochs):
            for xi, target in zip(X, y):
                pred = self.predict(xi)
                error = target - pred
                self.weights = [w + self.lr * error * x for w, x in zip(self.weights, xi)]
                self.bias += self.lr * error

# AND Gate example
X = [[0,0],[0,1],[1,0],[1,1]]
y = [0,0,0,1]
p = Perceptron(n_inputs=2)
p.train(X, y)
for xi in X:
    print(xi, "->", p.predict(xi))
```

NumPy Implementation

```
import numpy as np

class PerceptronNumpy:
    def __init__(self, n_inputs, lr=0.1, epochs=50):
        self.weights = np.zeros(n_inputs)
        self.bias = 0.0
        self.lr = lr
        self.epochs = epochs
    def activation(self, z):
```

```

        return np.where(z >= 0, 1, 0)

def predict(self, X):
    z = np.dot(X, self.weights) + self.bias
    return self.activation(z)

def train(self, X, y):
    for _ in range(self.epochs):
        for xi, target in zip(X, y):
            pred = self.predict(xi.reshape(1, -1))[0]
            error = target - pred
            self.weights += self.lr * error * xi
            self.bias += self.lr * error

X = np.array([[0,0],[0,1],[1,0],[1,1]])
y = np.array([0,0,0,1])
model = PerceptronNumpy(n_inputs=2)
model.train(X, y)
print(model.predict(X))

```

TensorFlow Implementation

```

import tensorflow as tf
import numpy as np

X = np.array([[0,0],[0,1],[1,0],[1,1]], dtype=np.float32)
y = np.array([[0],[0],[0],[1]], dtype=np.float32)

model = tf.keras.Sequential([
    tf.keras.layers.Dense(1, input_shape=(2,), activation='sigmoid')
])

model.compile(optimizer='sgd', loss='binary_crossentropy', metrics=['accuracy'])
model.fit(X, y, epochs=200, verbose=0)

print(model.predict(X))

```

2. NEURAL NETWORK

Simple Definition

A Neural Network is a collection of connected perceptron-like units (neurons) arranged in layers that work together to learn complex patterns from data, inspired loosely by the human brain.

Technical Definition

A Neural Network is a computational model composed of layers of interconnected nodes (neurons), where each connection has an associated weight. Information flows from an input layer through one or more hidden layers to an output layer, with non-linear activation functions enabling the network to approximate complex, non-linear functions through a process called forward propagation and backpropagation-based learning.

Hinglish Explanation

Neural Network basically bahut saare perceptrons ka combination hai jo layers mein arranged hote hain — input layer, hidden layer(s), aur output layer. Har neuron previous layer ke outputs ko input leta hai, weighted sum nikalta hai, activation function apply karta hai, aur next layer ko bhejta hai. Training ke time, network apna error nikalta hai aur backpropagation se weights update karta hai taaki agli baar prediction better ho.

Real-life Example


Predicting house prices: inputs like area, location, number of rooms, age of house go into the input layer; hidden layers learn complex interactions (e.g., "large area + good location = higher price"); output layer gives the predicted price. This non-linear modeling is something simple linear models often can't capture well.

Why it is used

- Can model complex, non-linear relationships that a single perceptron cannot.
- Solves problems like XOR that perceptrons cannot.
- Forms the basis for almost all deep learning architectures (ANN, CNN, RNN, Transformers).
- Highly flexible — used in classification, regression, generation tasks.

Architecture Diagram

Neural Network Architecture (Multi-Layer)



Step-by-step Working

- **Forward Propagation:** Input data passes through each layer; each neuron computes weighted sum + bias, then applies activation function (ReLU, sigmoid, tanh).
- **Loss Calculation:** Compare predicted output with actual label using a loss function (MSE for regression, Cross-Entropy for classification).
- **Backward Propagation:** Compute gradient of loss with respect to each weight using the chain rule, propagating error backward from output to input layer.
- **Weight Update:** Use an optimizer (SGD, Adam) to update weights in the direction that minimizes loss: $w = w - lr * \text{gradient}$.
- **Repeat** steps 1-4 for many epochs over the training dataset until loss converges.

Mathematical Formula with Explanation

Forward pass per neuron:

$$z = (W \cdot x) + b$$
$$a = \text{activation}(z) \quad [\text{e.g., } \text{ReLU}(z) = \max(0, z), \text{ Sigmoid}(z) = 1/(1+e^{-z})]$$

Loss Function (example - Mean Squared Error):

$$L = (1/n) * \sum (y_{\text{actual}} - y_{\text{predicted}})^2$$

Backpropagation (gradient via chain rule):

$$dL/dW = dL/da * da/dz * dz/dW$$

Weight Update (Gradient Descent):

$$w_{\text{new}} = w_{\text{old}} - \text{learning_rate} * dL/dW$$

Explanation: Forward pass computes predictions; loss measures error; backpropagation calculates how much each weight contributed to the error using calculus chain rule; gradient descent nudges weights to reduce that error iteratively.

Advantages & Disadvantages

Advantages:

- Can approximate any continuous function (Universal Approximation Theorem).
- Learns complex non-linear patterns automatically.
- Highly flexible for various tasks (classification, regression, generation).
- Forms foundation for advanced architectures.

Disadvantages:

- Requires large amounts of data to train well.
- Computationally expensive, needs GPUs for large networks.
- Prone to overfitting without regularization.
- Acts as a "black box" — hard to interpret decisions.
- Sensitive to hyperparameter choices (learning rate, layers, neurons).

Interview Questions

- What is forward propagation and backpropagation?
- Why do we need non-linear activation functions in a neural network?
- What is the vanishing gradient problem?
- Explain gradient descent and its variants (SGD, Adam, RMSProp).
- What is the Universal Approximation Theorem?
- How does a neural network differ from a perceptron?
- What is overfitting and how do you prevent it in neural networks?

Python Code

```
import random, math

def sigmoid(x):
    return 1 / (1 + math.exp(-x))

def sigmoid_derivative(x):
    return x * (1 - x)

class SimpleNeuralNetwork:
    def __init__(self, n_input, n_hidden, n_output):
        self.w1 = [[random.uniform(-1,1) for _ in range(n_hidden)] for _ in range(n_input)]
        self.w2 = [[random.uniform(-1,1) for _ in range(n_output)] for _ in range(n_hidden)]

    def forward(self, x):
        self.hidden = [sigmoid(sum(x[i]*self.w1[i][j] for i in range(len(x)))) for j in range(len(self.w1[0]))]
        self.output = [sigmoid(sum(self.hidden[i]*self.w2[i][j] for i in range(len(self.hidden)))) for j in range(len(self.w2[0]))]
        return self.output

nn = SimpleNeuralNetwork(2, 3, 1)
print(nn.forward([0.5, 0.8]))
```

NumPy Implementation

```
import numpy as np

def sigmoid(x): return 1 / (1 + np.exp(-x))
def sigmoid_deriv(x): return x * (1 - x)

X = np.array([[0,0],[0,1],[1,0],[1,1]])
y = np.array([[0],[1],[1],[0]]) # XOR problem

np.random.seed(1)
w1 = np.random.uniform(-1,1,(2,4))
```

```

W2 = np.random.uniform(-1,1,(4,1))
lr = 0.5

for epoch in range(10000):
    # Forward
    hidden = sigmoid(np.dot(X, W1))
    output = sigmoid(np.dot(hidden, W2))

    # Backward
    error = y - output
    d_output = error * sigmoid_deriv(output)
    error_hidden = d_output.dot(W2.T)
    d_hidden = error_hidden * sigmoid_deriv(hidden)

    W2 += hidden.T.dot(d_output) * lr
    W1 += X.T.dot(d_hidden) * lr

print(output)

```

TensorFlow Implementation

```

import tensorflow as tf
import numpy as np

X = np.array([[0,0],[0,1],[1,0],[1,1]], dtype=np.float32)
y = np.array([[0],[1],[1],[0]], dtype=np.float32) # XOR

model = tf.keras.Sequential([
    tf.keras.layers.Dense(4, activation='relu', input_shape=(2,)),
    tf.keras.layers.Dense(1, activation='sigmoid')
])

model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
model.fit(X, y, epochs=500, verbose=0)
print(model.predict(X))

```

3. ARTIFICIAL NEURAL NETWORK (ANN)

Simple Definition

An Artificial Neural Network (ANN) is a fully-connected multi-layer neural network — also called a Multi-Layer Perceptron (MLP) — designed to learn from structured/tabular data through layers of densely connected neurons.

Technical Definition

ANN is a feedforward neural network architecture consisting of an input layer, one or more fully connected (dense) hidden layers, and an output layer, where every neuron in one layer is connected to every neuron in the next layer. It uses non-linear activation functions and is trained via backpropagation and gradient-based optimization to minimize a loss function, commonly applied to structured/tabular data classification and regression tasks.

Hinglish Explanation

ANN basically ek "fully connected" neural network hota hai, jisme har layer ka har neuron agle layer ke har neuron se connected hota hai (isiliye "Dense" layers bolte hain). Yeh general purpose network hai jo tabular data (jaise Excel sheets, structured datasets) pe achha kaam karta hai — jaise customer churn prediction, loan approval, ya price prediction. CNN/RNN se iska difference yeh hai ki ANN mein koi special structure (jaise spatial ya sequential) assume nahi kiya jata.


Real-life Example

A telecom company predicting customer churn (will a customer leave or stay) using features like monthly bill, contract type, number of complaints, tenure, and data usage. An ANN takes these tabular features as input and outputs a probability of churn — used widely in banking, insurance, and HR analytics for tabular prediction tasks.

Why it is used

- Excellent for structured/tabular data where there's no spatial or sequential pattern.
- Can model complex, non-linear relationships between many features.
- Flexible architecture — can be customized with any number of layers/neurons.
- Used as the final classification/regression "head" even within CNNs and RNNs.

Architecture Diagram



Step-by-step Working

- Input features (numerical/encoded categorical) are fed into the input layer (no computation here, just passes values).
- Each hidden layer neuron computes: $z = (w \cdot x) + b$, then applies activation (ReLU is most common).
- Output is passed to the next dense layer, repeating the same process layer by layer.
- Final output layer applies an appropriate activation: Sigmoid (binary classification), Softmax (multi-class), or Linear (regression).
- Loss is computed comparing prediction vs actual label.
- Backpropagation computes gradients of loss w.r.t. all weights across all layers using chain rule.
- Optimizer (Adam/SGD) updates weights to minimize loss.
- Process repeats across many epochs with techniques like dropout/batch norm to prevent overfitting.

Mathematical Formula with Explanation

Layer-wise computation (for layer l):

$$Z[l] = W[l] \cdot A[l-1] + b[l]$$

$$A[l] = \text{activation}(Z[l])$$

Softmax (multi-class output layer):

$$\text{softmax}(z_i) = e^{(z_i)} / \sum(e^{(z_j)}) \text{ for all } j$$

Cross-Entropy Loss (classification):

$$L = -\sum(y_{\text{actual}} * \log(y_{\text{predicted}}))$$

Adam Optimizer Update (simplified):

$$m = \beta_1 * m + (1 - \beta_1) * \text{grad}$$

$$v = \beta_2 * v + (1 - \beta_2) * \text{grad}^2$$

$$W = W - \text{lr} * m / (\sqrt{v} + \epsilon)$$

Explanation: $Z[l]$ and $A[l]$ represent pre-activation and post-activation values per layer. Softmax converts raw scores into probabilities summing to 1. Cross-entropy penalizes wrong confident predictions heavily.

Adam combines momentum and adaptive learning rates for faster, stable convergence than plain gradient descent.

Advantages & Disadvantages

Advantages:

- Works very well on structured/tabular data.
- Simple, flexible architecture, easy to design and scale.
- Can approximate complex non-linear functions (Universal Approximation).
- Well-supported by all major frameworks (TensorFlow, PyTorch).

Disadvantages:

- Doesn't exploit spatial structure (bad for images) or sequence order (bad for text/time-series) — needs CNN/RNN for those.
- Fully connected layers have many parameters → prone to overfitting and slow training on large inputs.
- Requires significant feature engineering for tabular data sometimes.
- Hard to interpret ("black box") compared to simpler models like decision trees.

Interview Questions

- What is the difference between ANN and a simple Neural Network?
- Why is ANN not ideal for image data compared to CNN?
- What is a Dense/Fully-Connected layer?
- Explain the role of dropout and batch normalization in ANN.
- What is the difference between Sigmoid, Softmax, and ReLU activation functions?
- How do you decide the number of hidden layers and neurons?
- What is overfitting in ANN and how do you handle it (regularization, dropout, early stopping)?

Python Code

```
import numpy as np

def relu(x): return np.maximum(0, x)
def relu_deriv(x): return (x > 0).astype(float)
def sigmoid(x): return 1 / (1 + np.exp(-x))

class ANN:
    def __init__(self, layer_sizes):
        self.weights = [np.random.randn(layer_sizes[i], layer_sizes[i+1]) * 0.1
                        for i in range(len(layer_sizes)-1)]
        self.biases = [np.zeros((1, layer_sizes[i+1])) for i in range(len(layer_sizes)-1)]

    def forward(self, X):
        a = X
        for i in range(len(self.weights)-1):
            a = relu(np.dot(a, self.weights[i]) + self.biases[i])
        out = sigmoid(np.dot(a, self.weights[-1]) + self.biases[-1])
        return out

model = ANN([4, 8, 8, 1])
X = np.random.randn(5, 4)
print(model.forward(X))
```

NumPy Implementation

```
import numpy as np

# Full ANN with forward + backward pass (binary classification)
class ANNNumpy:
```

```

def __init__(self, n_input, n_hidden, n_output, lr=0.1):
    self.W1 = np.random.randn(n_input, n_hidden) * 0.1
    self.b1 = np.zeros((1, n_hidden))
    self.W2 = np.random.randn(n_hidden, n_output) * 0.1
    self.b2 = np.zeros((1, n_output))
    self.lr = lr

def relu(self, x): return np.maximum(0, x)
def relu_deriv(self, x): return (x > 0).astype(float)
def sigmoid(self, x): return 1/(1+np.exp(-x))

def forward(self, X):
    self.Z1 = np.dot(X, self.W1) + self.b1
    self.A1 = self.relu(self.Z1)
    self.Z2 = np.dot(self.A1, self.W2) + self.b2
    self.A2 = self.sigmoid(self.Z2)
    return self.A2

def backward(self, X, y):
    m = X.shape[0]
    dZ2 = self.A2 - y
    dW2 = np.dot(self.A1.T, dZ2) / m
    db2 = np.sum(dZ2, axis=0, keepdims=True) / m
    dA1 = np.dot(dZ2, self.W2.T)
    dZ1 = dA1 * self.relu_deriv(self.Z1)
    dW1 = np.dot(X.T, dZ1) / m
    db1 = np.sum(dZ1, axis=0, keepdims=True) / m

    self.W2 -= self.lr * dW2; self.b2 -= self.lr * db2
    self.W1 -= self.lr * dW1; self.b1 -= self.lr * db1

X = np.random.randn(100, 4)
y = (np.sum(X, axis=1, keepdims=True) > 0).astype(float)
model = ANNNumpy(4, 8, 1)
for epoch in range(1000):
    pred = model.forward(X)
    model.backward(X, y)
print("Final prediction sample:", model.forward(X[:5]))

```

TensorFlow Implementation

```

import tensorflow as tf
from sklearn.datasets import make_classification

X, y = make_classification(n_samples=500, n_features=10, random_state=42)

model = tf.keras.Sequential([
    tf.keras.layers.Dense(32, activation='relu', input_shape=(10,)),
    tf.keras.layers.Dropout(0.3),
    tf.keras.layers.Dense(16, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])

model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
model.fit(X, y, epochs=20, batch_size=16, validation_split=0.2, verbose=0)
print(model.evaluate(X, y))

```

4. CONVOLUTIONAL NEURAL NETWORK (CNN)

Simple Definition

A Convolutional Neural Network (CNN) is a specialized neural network designed to process grid-like data such as images, by automatically learning spatial features (edges, shapes, textures, objects) using small sliding filters called kernels.

Technical Definition

CNN is a deep learning architecture that applies learnable convolutional filters (kernels) across input data to extract local spatial hierarchies of features, typically followed by pooling layers for dimensionality reduction and fully connected layers for final classification/regression. It exploits parameter sharing and local connectivity to drastically reduce the number of parameters compared to a fully connected ANN, making it highly effective for image, video, and spatial-grid data.

Hinglish Explanation

CNN ek special type ka neural network hai jo images jaise grid-data ko samajhne ke liye design kiya gaya hai. Yeh chhote-chhote filters (kernels) ko image par slide karta hai taaki edges, corners, textures, aur patterns detect kar sake. Pehle layers simple features (edges) detect karte hain, deeper layers complex features (eyes, face, object shapes) detect karte hain. Pooling layers image size ko chhota karke important information retain karte hain, aur end mein fully connected layers final classification karte hain (jaise "yeh cat hai ya dog").

Real-life Example


Face unlock feature in smartphones: CNN takes the camera image as input, early layers detect edges and contours, middle layers detect features like eyes/nose/mouth shapes, deeper layers combine them to recognize the unique facial pattern, and the final layer matches it against the stored face — used in apps like Face ID, medical image diagnosis (X-ray/MRI classification), and self-driving car object detection.

Why it is used

- Automatically learns spatial hierarchies of features without manual feature engineering.
- Parameter sharing (same filter reused across the image) drastically reduces parameters vs ANN.
- Translation invariance — can detect an object regardless of its position in the image.
- State-of-the-art performance for image classification, object detection, segmentation, and video analysis.

Architecture Diagram

Convolutional Neural Network (CNN) Architecture



Step-by-step Working

- **Convolution Operation:** A small filter/kernel (e.g., 3x3) slides over the input image, computing element-wise multiplication + sum at each position to produce a "feature map" that highlights patterns like edges.
- **Activation (ReLU):** Apply ReLU to introduce non-linearity, keeping only positive activations.
- **Pooling (Max/Average):** Reduce spatial dimensions (e.g., 2x2 max pooling picks the largest value in each region), reducing computation and adding translation invariance.
- **Stacking Layers:** Repeat Conv → ReLU → Pool multiple times; early layers learn simple features (edges, colors), deeper layers learn complex features (shapes, objects).
- **Flatten:** Convert the final 2D/3D feature maps into a 1D vector.
- **Fully Connected (Dense) Layers:** Combine extracted features to make final decision.
- **Output Layer:** Softmax (multi-class) or Sigmoid (binary) gives final class probabilities.
- **Backpropagation:** Loss is computed and gradients flow back through FC layers, pooling, and conv filters, updating filter weights via gradient descent.

Mathematical Formula with Explanation

Convolution Operation:

$$S(i,j) = \sum_m \sum_n I(i+m, j+n) * K(m,n)$$

where I is the input image, K is the kernel/filter, and S is the resulting feature map.

Output size formula:

$$\text{Output_size} = ((\text{Input_size} - \text{Filter_size} + 2 * \text{Padding}) / \text{Stride}) + 1$$

Max Pooling:

$$P(i,j) = \max(\text{region of feature map covered by pooling window})$$

ReLU Activation:

$$\text{ReLU}(x) = \max(0, x)$$

Explanation: Convolution slides the kernel over the image performing dot products to detect local patterns. The output-size formula tells us the resulting feature map dimensions based on padding and stride choices. Pooling compresses feature maps while retaining dominant features. ReLU adds non-linearity so the network can learn complex patterns, not just linear combinations.

Advantages & Disadvantages

Advantages:

- Excellent at capturing spatial/local patterns in images and video.
- Parameter sharing drastically reduces number of trainable parameters vs ANN.
- Translation invariant — detects objects regardless of position.
- Automatically learns hierarchical features (no manual feature engineering needed).
- State of the art for vision tasks (classification, detection, segmentation).

Disadvantages:

- Computationally expensive, needs large datasets and GPUs for training.
- Not naturally suited for non-grid/sequential data (text, time series) — needs RNN/Transformer instead.
- Requires careful tuning of architecture (filters, layers, strides).
- Can be prone to overfitting on small datasets.
- Less interpretable — hard to explain exactly what filters learn ("black box" deep model).

Interview Questions

- What is a convolution operation and how does it differ from matrix multiplication?
- Why is CNN better than ANN for image data?
- What is the purpose of pooling layers? Difference between Max Pooling and Average Pooling?
- Explain padding and stride and how they affect output size.
- What is parameter sharing and why is it important in CNN?
- What is the role of the Flatten layer in CNN architecture?
- Explain popular CNN architectures (LeNet, AlexNet, VGG, ResNet) briefly.
- What is the vanishing gradient problem in deep CNNs and how do architectures like ResNet solve it?

Python Code

```
import numpy as np

def conv2d(image, kernel, stride=1, padding=0):
    if padding > 0:
        image = np.pad(image, ((padding,padding),(padding,padding)), mode='constant')
    ih, iw = image.shape
    kh, kw = kernel.shape
    oh = (ih - kh)//stride + 1
    ow = (iw - kw)//stride + 1
    output = np.zeros((oh, ow))
    for i in range(0, oh):
        for j in range(0, ow):
            region = image[i*stride:i*stride+kh, j*stride:j*stride+kw]
            output[i, j] = np.sum(region * kernel)
    return output

# Example: edge detection kernel
image = np.random.rand(8, 8)
edge_kernel = np.array([[ -1, -1, -1], [-1, 8, -1], [-1, -1, -1]])
feature_map = conv2d(image, edge_kernel)
print(feature_map.shape)
print(feature_map)
```

NumPy Implementation

```
import numpy as np

def relu(x): return np.maximum(0, x)

def max_pool(feature_map, size=2, stride=2):
```

```

h, w = feature_map.shape
oh = (h - size)//stride + 1
ow = (w - size)//stride + 1
pooled = np.zeros((oh, ow))
for i in range(oh):
    for j in range(ow):
        region = feature_map[i*stride:i*stride+size, j*stride:j*stride+size]
        pooled[i, j] = np.max(region)
return pooled

def conv2d(image, kernel, stride=1, padding=0):
    if padding > 0:
        image = np.pad(image, ((padding,padding),(padding,padding)), mode='constant')
    ih, iw = image.shape
    kh, kw = kernel.shape
    oh = (ih - kh)//stride + 1
    ow = (iw - kw)//stride + 1
    output = np.zeros((oh, ow))
    for i in range(oh):
        for j in range(ow):
            region = image[i*stride:i*stride+kh, j*stride:j*stride+kw]
            output[i, j] = np.sum(region * kernel)
    return output

# Simple CNN forward pass: Conv -> ReLU -> Pool
image = np.random.rand(10, 10)
kernel = np.array([[1,0,-1],[1,0,-1],[1,0,-1]]) # vertical edge detector

conv_out = conv2d(image, kernel, stride=1, padding=0)
relu_out = relu(conv_out)
pooled_out = max_pool(relu_out, size=2, stride=2)

print("Conv output shape:", conv_out.shape)
print("Pooled output shape:", pooled_out.shape)

```

TensorFlow Implementation

```

import tensorflow as tf
from tensorflow.keras import layers, models

# CNN for image classification (e.g., CIFAR-10 style: 32x32x3 images, 10 classes)
model = models.Sequential([
    layers.Conv2D(32, (3,3), activation='relu', padding='same', input_shape=(32,32,3)),
    layers.MaxPooling2D((2,2)),
    layers.Conv2D(64, (3,3), activation='relu', padding='same'),
    layers.MaxPooling2D((2,2)),
    layers.Conv2D(128, (3,3), activation='relu', padding='same'),
    layers.MaxPooling2D((2,2)),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dropout(0.5),
    layers.Dense(10, activation='softmax')
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.summary()

# Example training (using CIFAR-10 dataset)
(X_train, y_train), (X_test, y_test) = tf.keras.datasets.cifar10.load_data()
X_train, X_test = X_train/255.0, X_test/255.0
model.fit(X_train, y_train, epochs=5, validation_data=(X_test, y_test))

```

Quick Comparison Table

Aspect	Perceptron	Neural Network	ANN	CNN
Layers	Single layer	Multiple layers	Multiple dense layers	Conv + Pool + Dense layers
Data type	Linearly separable	General	Tabular/structured	Images/grid data
Activation	Step function	Sigmoid/ReLU/Tanh	ReLU/Softmax/Sigmoid	ReLU + Softmax
Key strength	Simplicity	Non-linear learning	Tabular pattern learning	Spatial feature learning
Limitation	No XOR solving	Needs lots of data	No spatial awareness	Computationally heavy

End of Document — Deep Learning Fundamentals Guide